



DSC630

Week 8 solution

Abhishek Srivastava

8.2 Assignment: Time Series Modeling

You will be using the dataset us_retail_sales.csv for this assignment. This data gives the total monthly retail sales in the US from January 1992 until June 2021. With this dataset, complete the following steps: Plot the data with proper labeling and make some observations on the graph. Split this data into a training and test set. Use the last year of data (July 2020 – June 2021) of data as your test set and the rest as your training set. Use the training set to build a predictive model for the monthly retail sales. Use the model to predict the monthly retail sales on the last year of data. Report the RMSE of the model predictions on the test set.

```
In [23]: # Load required libraries
import pandas as pd
import numpy as np

# Load data file
df = pd.read_csv('us_retail_sales.csv')
# Display the first few rows of the dataframe
print(df.head())
# Display the column names
print("\nColumns in the dataframe:")
print(df.columns)
```

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	\
0	1992	146925	147223	146805	148032	149010	149800	150761.0	151067.0	
1	1993	157555	156266	154752	158979	160605	160127	162816.0	162506.0	
2	1994	167518	169649	172766	173106	172329	174241	174781.0	177295.0	
3	1995	182413	179488	181013	181686	183536	186081	185431.0	186806.0	
4	1996	189135	192266	194029	194744	196205	196136	196187.0	196218.0	

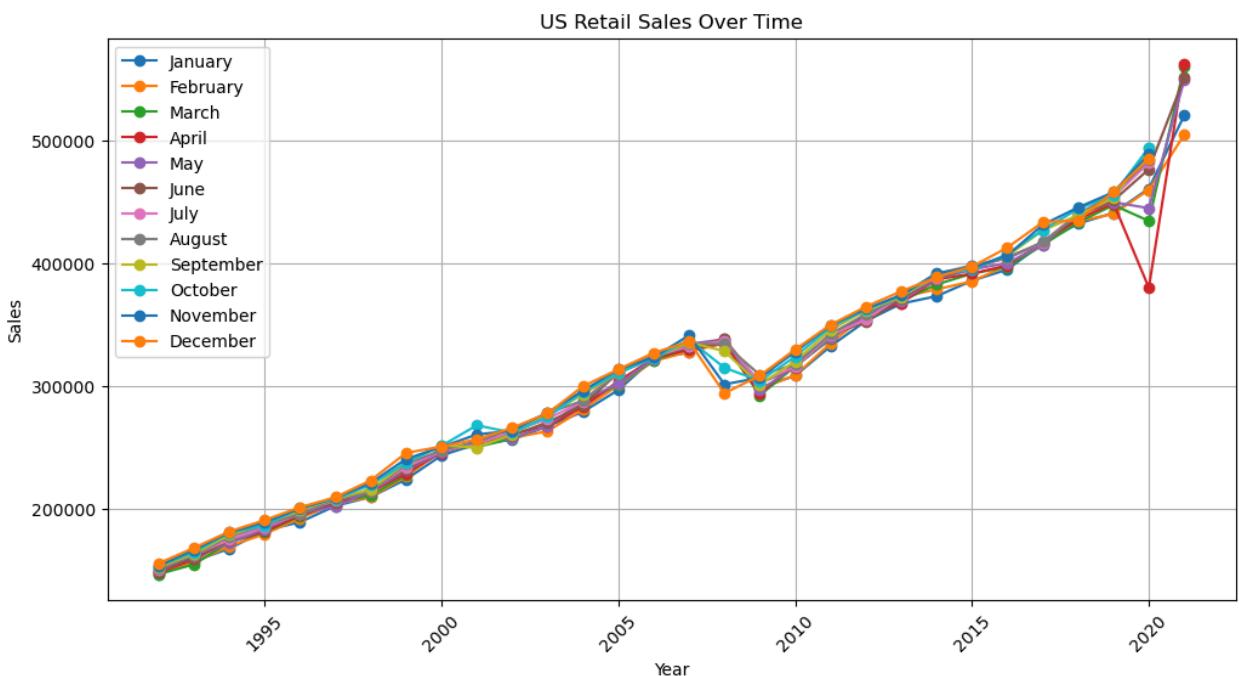
	SEP	OCT	NOV	DEC
0	152588.0	153521.0	153583.0	155614.0
1	163258.0	164685.0	166594.0	168161.0
2	178787.0	180561.0	180703.0	181524.0
3	187366.0	186565.0	189055.0	190774.0
4	198859.0	200509.0	200174.0	201284.0

Columns in the dataframe:

```
Index(['YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL', 'AUG', 'SEP',
       'OCT', 'NOV', 'DEC'],
      dtype='object')
```

1. Plot the data with proper labeling and make some observations on the graph.

```
In [24]: # Plot the data with proper labeling and make some observations on the graph w  
# Columns in the dataframe:  
#Index(['YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL', 'AUG', 'SEP',  
#       'OCT', 'NOV', 'DEC'],  
#      dtype='object')  
import matplotlib.pyplot as plt  
plt.figure(figsize=(12, 6))  
plt.plot(df['YEAR'], df['JAN'], marker='o', label='January')  
plt.plot(df['YEAR'], df['FEB'], marker='o', label='February')  
plt.plot(df['YEAR'], df['MAR'], marker='o', label='March')  
plt.plot(df['YEAR'], df['APR'], marker='o', label='April')  
plt.plot(df['YEAR'], df['MAY'], marker='o', label='May')  
plt.plot(df['YEAR'], df['JUN'], marker='o', label='June')  
plt.plot(df['YEAR'], df['JUL'], marker='o', label='July')  
plt.plot(df['YEAR'], df['AUG'], marker='o', label='August')  
plt.plot(df['YEAR'], df['SEP'], marker='o', label='September')  
plt.plot(df['YEAR'], df['OCT'], marker='o', label='October')  
plt.plot(df['YEAR'], df['NOV'], marker='o', label='November')  
plt.plot(df['YEAR'], df['DEC'], marker='o', label='December')  
plt.title('US Retail Sales Over Time')  
plt.xlabel('Year')  
plt.ylabel('Sales')  
plt.xticks(rotation=45)  
plt.grid()  
plt.legend()  
plt.show()
```



2. Some observations on the graph

- US retail sales mostly go up from the early 1990s to 2007. Sales get bigger every year for all months.
- There is a clear drop in sales around 2008 and 2009. This happened because of the global financial crisis, which hurt the economy everywhere.
- After that, sales start growing again and keep rising until about 2019. The economy gets better and people spend more.
- There is a very large drop in 2020, mostly in March and April. This lines up with the start of the COVID-19 pandemic, when stores closed and people stayed home.
- Sales quickly go back up at the end of 2020 and in 2021, reaching the highest levels ever. This happens as stores reopen and the government helps people with stimulus money.
- Sales are highest every December, which is the holiday season when people buy a lot. Sales are lowest in January and February, just after the holidays.
- Even though there are a few drops, US retail sales mostly keep rising over time. This shows that retail business in the US is strong and can recover from problems.

3. Split this data into a training and test set. Use the last year of data (July 2020 – June 2021) of data as your test set and the rest as your training set.

```
In [25]: # Split this data into a training and test set. Use the last year of data (Jul
# of data as your test set and the rest as your training set.
# # Columns in the dataframe:
#Index(['YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL', 'AUG', 'SEP',
#       'OCT', 'NOV', 'DEC'],
#      dtype='object')

# First, we need to reshape the dataframe to have a 'MONTH' column
df = df.melt(id_vars=['YEAR'], var_name='MONTH', value_name='SALES')
# Create a 'YYYYMM' column for easier filtering
# Cast value to numeric for df['YYYYMM']
df['YYYYMM'] = df['YEAR'].astype(str) + df['MONTH'].map({
    'JAN': '01', 'FEB': '02', 'MAR': '03', 'APR': '04', 'MAY': '05', 'JUN': '06',
    'JUL': '07', 'AUG': '08', 'SEP': '09', 'OCT': '10', 'NOV': '11', 'DEC': '12'})
df['YYYYMM'] = pd.to_numeric(df['YYYYMM'])
# Define the test set period
test_start = 202007
```

```
test_end = 202106
# Split the data
train_df = df[(df['YYYYMM'] < test_start)]
test_df = df[(df['YYYYMM'] >= test_start) & (df['YYYYMM'] <= test_end)]
```

4. Use the training set to build a predictive model for the monthly retail sales

```
In [26]: # Use whole training set to build a predictive model for the monthly retail sa
from sklearn.linear_model import LinearRegression
# Prepare the data for modeling
X_train = train_df[['YYYYMM']]
y_train = train_df['SALES']
# Build the model
model = LinearRegression()
model.fit(X_train, y_train)
# Display model coefficients
print("Model Coefficients:")
print(f"Intercept: {model.intercept_}")
print(f"Coefficient: {model.coef_}")
```

```
Model Coefficients:
Intercept: -20697841.76585196
Coefficient: [104.68329077]
```

5. Use the model to predict the monthly retail sales on the last year of data

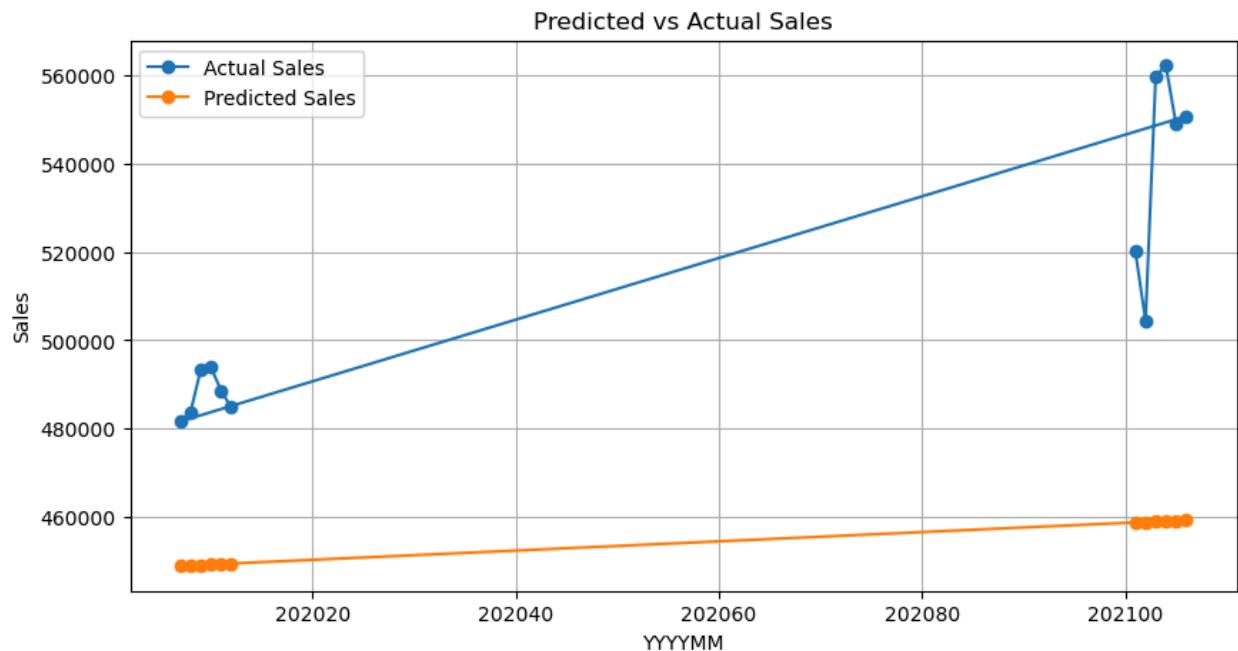
```
In [27]: # Prepare the test labels (actual sales)
y_test = test_df['SALES']
X_test = test_df[['YYYYMM']]
# Predict the sales on the test set
y_pred = model.predict(X_test)

# Display the predicted vs actual sales for each test month
print("\nPredicted vs Actual Sales:")
for i in range(len(X_test)):
    print(f"Month: {X_test.iloc[i]['YYYYMM']}, Predicted Sales: {y_pred[i]:.2f}")

# Optionally plot predicted vs actual for visual comparison
import matplotlib.pyplot as plt
plt.figure(figsize=(10, 5))
plt.plot(X_test['YYYYMM'], y_test, marker='o', label='Actual Sales')
plt.plot(X_test['YYYYMM'], y_pred, marker='o', label='Predicted Sales')
plt.xlabel('YYYYMM')
plt.ylabel('Sales')
plt.title('Predicted vs Actual Sales')
plt.legend()
plt.grid()
plt.show()
```

Predicted vs Actual Sales:

```
Month: 202101, Predicted Sales: 458755.98, Actual Sales: 520162.00
Month: 202102, Predicted Sales: 458860.67, Actual Sales: 504458.00
Month: 202103, Predicted Sales: 458965.35, Actual Sales: 559871.00
Month: 202104, Predicted Sales: 459070.03, Actual Sales: 562269.00
Month: 202105, Predicted Sales: 459174.72, Actual Sales: 548987.00
Month: 202106, Predicted Sales: 459279.40, Actual Sales: 550782.00
Month: 202007, Predicted Sales: 448915.75, Actual Sales: 481627.00
Month: 202008, Predicted Sales: 449020.44, Actual Sales: 483716.00
Month: 202009, Predicted Sales: 449125.12, Actual Sales: 493327.00
Month: 202010, Predicted Sales: 449229.80, Actual Sales: 493991.00
Month: 202011, Predicted Sales: 449334.49, Actual Sales: 488652.00
Month: 202012, Predicted Sales: 449439.17, Actual Sales: 484782.00
```



```
In [28]: # Report the RMSE of the model predictions on the test set.
```

```
from sklearn.metrics import mean_squared_error
rmse = np.sqrt(mean_squared_error(test_df['SALES'], y_pred))
print(f"\nRoot Mean Squared Error (RMSE) on Test Set: {rmse:.2f}")
```

Root Mean Squared Error (RMSE) on Test Set: 65920.99